

Comparative Analysis of Fault Tolerance in Elixir and Other Distributed Languages

Nuno Ribeiro

**A dissertation submitted in partial fulfillment of
the requirements for the degree of Master of Science,
Specialisation Area of Software Engineering**

Advisor: Dr. Luís Nogueira

Porto, December 11, 2024

Statement of Integrity

I hereby declare having conducted this academic work with integrity.

I have not plagiarised or applied any form of undue use of information or falsification of results along the process leading to its elaboration.

Therefore the work presented in this document is original and authored by me, having not previously been used for any other end.

I further declare that I have fully acknowledged the Code of Ethical Conduct of P.PORTO.

ISEP, Porto, December 11, 2024

Contents

List of Figures	vii
List of Tables	ix
1 Introduction	1
1.1 Context and Problem	1
1.2 Objectives	1
1.3 Ethical Considerations	1
1.4 Document structure	2
2 Background	3
2.1 Distributed Systems	3
2.1.1 Characteristics	3
Transparency	3
Reliability and Availability	3
Scalability	4
Fault tolerance	4
2.1.2 Communication	4
2.1.3 Challenges	4
CAP theorem	4
2.2 Fault Tolerance	5
2.2.1 Fault Tolerance Taxonomy	5
2.2.2 Strategies	6
Retry Mechanism	6
Circuit Breaker Pattern	6
Replication and Redundancy	7
Checkpointing and Message Logging	11
2.3 Distributed and Concurrency Programming	11
2.3.1 Models and Paradigms	12
Actor Model	12
Communicating Sequential Processes	12
Microservices Architectures	13
2.3.2 Distributed and Concurrent Programming Languages	13
Analyses and Language Choice Justification	14
3 Literature Review	17
3.1 Research Questions	17
3.2 State of Art	17
3.3 Conclusions	17
4 Planning	19

4.1	Project Charter	19
4.2	Work Breakdown Structure	20
4.3	Gantt Diagram	23
4.3.1	Project Management and Scheduling	23
4.3.2	Monitoring and Controlling Procedures	23
4.3.3	Meeting Sessions	24
4.3.4	Competencies Development Plan	24
4.4	Risk Management	26
4.4.1	Risk 1: Bugs in Third-Party Libraries	26
4.4.2	Risk 2: Integration Challenges Between Components	27

List of Figures

2.1	Circuit Breaker States	7
2.2	State Diagram of Nodes in Raft	10
4.1	The Work Breakdown Structure (WBS) of the project.	20
4.2	Monitoring and control procedures displayed on the Gantt chart.	24
4.3	Meeting sessions represented on the Gantt chart.	25
4.4	Competencies Development Plan represented on the Gantt chart.	25

List of Tables

2.1	Brief Description of Failure Types	5
2.2	Characteristics of Distributed and Concurrent Programming Languages . .	15
4.1	WBS dictionary	23

Chapter 1

Introduction

1.1 Context and Problem

- Lack of information focusing on fault tolerance comparison
- Old information available
- Increase usage of microservices / distributed systems

1.2 Objectives

1.3 Ethical Considerations

Ethical considerations play a crucial role in software engineering research, ensuring the integrity, transparency, and societal relevance of the work. This section outlines the ethical principles applied.

Transparency and Fairness in Results. As this research is focused on the comparison of programming languages, it is essential to maintain impartiality and avoid any bias in the results. Research integrity demands that results are not manipulated or altered to provoke more appealing discussions or gain community approval [1]. This dissertation adheres to the principle of transparency, ensuring that the benchmarking results reflect the true performance of each language.

Replicability and Verification. Replication is an important aspect of scientific research, enabling others to validate findings [2]. This dissertation involves the development of prototypes and proof-of-concepts to evaluate specific fault-tolerance strategies. To uphold ethical standards, all tests and configurations will be documented on a public repository to allow replication.

Adherence to Professional Codes of Ethics. This work adheres to the ethical principles outlined in the Institute of Electrical and Electronics Engineers (IEEE) Code of Ethics [3] and the Association for Computing Machinery Code (ACM) of Ethics and Professional Conduct [4], which emphasize integrity, respect, fairness, and authorized use of content. Researchers must act responsibly by avoiding any practices that could harm the reputation or fairness of the comparison, maintaining an ethical commitment to the broader community of developers and researchers.

Avoidance of Plagiarism and Proper Citation. Plagiarism undermines the credibility and value of academic work. In alignment with the Code of Good Practices and Conduct of

Polytechnic of Porto, particularly Article 10, this dissertation ensures proper attribution of all referenced works. Accurate citation is fundamental to acknowledge the contributions of others, demonstrate the research's academic honesty, and respect intellectual property.

1.4 Document structure

Chapter 2

Background

2.1 Distributed Systems

In the early days of computing, computers were large and expensive, operating as standalone machines without the ability to communicate with each other. As technology advanced, smaller and more affordable computers, such as smartphones and other devices, were developed, along with high-speed networking that allowed connectivity across a network [5]. These innovations made it possible to create systems distributed across nodes where tasks could be processed collectively to achieve a common goal [5]. Nodes in a distributed system may refer to physical devices or software processes [6].

To the end-user, distributed systems appear as a single, large virtual system, making the underlying logic transparent [6]. These systems achieve a shared objective by transmitting messages through various nodes and dividing computational tasks among them, increasing resilience and isolating business logic [6, 7]. Distributed systems can present heterogeneity, such as differing clocks, memory, programming languages, operating systems, or geographical locations, all of which must be abstracted from the end-user [5, 7].

2.1.1 Characteristics

On a distributed system, when being well-structured, it is possible to find, among others, the following most popular characteristic:

Transparency

Transparency in distributed systems enables seamless user interaction by hiding the complexity of underlying operations [5, 8]. Key aspects include access transparency, which allows resource usage without concern for system differences, and location transparency, which hides the physical location of resources, as seen with Uniform Resource Locators (URLs) [5, 9]. Replication transparency ensures reliability by masking data duplication, while failure transparency enables systems to handle faults without user disruption [5, 9]. Together, these forms of transparency enhance usability, robustness, and reliability.

Reliability and Availability

A distributed system should have reliability and availability aspects. Reliability refers to its ability to continuously perform its intended requirements without interruption, operating exactly as designed, even in the presence of certain internal failures [10]. A highly reliable

system maintains consistent, uninterrupted service over an extended period, minimizing disruptions for users [5], on other hand, availability measures the probability that the system is operational and ready to respond correctly at any given moment, often expressed as a percentage of system up-time [5, 11].

Scalability

Designing and building a distributed system is complex, but also enables the creation of highly scalable systems, capable of expanding to meet increasing demands [5, 6, 12]. This characteristic is particularly evident as cloud-based systems become more popular, allowing users to interact with applications over the internet rather than relying on local desktop computing power [13]. Cloud services must support a large volume of simultaneous connections and interactions, making scalability a crucial factor [5].

Fault tolerance

Fault tolerance is a critical characteristic of distributed systems, closely linked to reliability, availability, and scalability. For a system to maintain these properties, it must be able to mask failures and continue operating despite the presence of errors [5]. Fault tolerance is especially vital in distributed environments where system failures can lead to significant disruptions and economic losses across sectors such as finance, telecommunications, and transportation [7].

The primary goal of a fault-tolerant system is to enable continuous operation by employing specific strategies and design patterns to mask the possible errors [14].

2.1.2 Communication

Communication is fundamental in distributed systems for coordination and data exchange. Nodes communicate over networks or via Interprocess Communication (IPC) when on the same machine [6]. Synchronous communication involves blocking operations where the sender waits for a response, suitable for scenarios requiring confirmation [5, 9]. In contrast, asynchronous communication allows non-blocking operations, enabling the sender to proceed without waiting. This approach, often supported by message queues, is ideal for decoupled and heterogeneous systems [5].

2.1.3 Challenges

Distributed systems encounter numerous challenges, including scalability [10], managing software, network, and disk failures [15, 16], heterogeneity [9], coordination among nodes [6], and difficulties on debugging and testing [16, 17]. For the scope of this dissertation only the CAP theorem will be discussed.

CAP theorem

The CAP theorem says that in a system where nodes are networked and share data, it is impossible to simultaneously achieve all three properties of Consistency, Availability, and Partition Tolerance [5, 6]. This theorem underlines a critical trade-off in distributed systems: only two of these properties can be fully ensured at any given time [18, 19]. A description of the properties can be given by:

- **Consistency:** Ensures that all nodes in the system reflect the same data at any time, so each read returns the latest write.
- **Availability:** Guarantees that every request receives a response, whether successful or not, even if some nodes are offline.
- **Partition tolerance:** Allows the system to continue operating despite network partitions, where nodes may temporarily lose the ability to communicate.

According to the CAP theorem, when a network partition occurs, a distributed system must prioritize either consistency or availability, as achieving all three properties is not feasible in practice [5, 6, 18]. This concept is directly relevant to this dissertation, as fault tolerance strategies discussed later will account for these trade-offs to optimize specific properties.

2.2 Fault Tolerance

With the extensive use of software systems across various domains, the demand for reliable and available systems is essential. However, errors in software are inevitable, making fault tolerance a critical attribute for systems to continue functioning correctly even in the presence of failures [7]. Fault tolerance can address a range of issues, including networking, hardware, software, and other dimensions, with various strategies designed to manage these different fault types [5, 20].

2.2.1 Fault Tolerance Taxonomy

It is important to classify and understand the types of failures that can arise. This section presents a taxonomy of fault tolerance concepts, drawing on the framework proposed by Isukapalli et al.[21]. A fault is defined as an underlying defect within a system component that may lead to an error, which is a deviation from the intended internal state. If this error remains unresolved, it may escalate into a system failure, potentially impacting system functionality either partially or completely [21, 22].

Failures are the external manifestations of the internal faults, as outlined in Table 2.1. These include crash failures, where the system halts entirely, to arbitrary failures, where responses are erratic and potentially misleading [5].

Type of Failure	Description
Crash Failure	The system halts and stops all operations entirely. Although it was functioning correctly before the halt, it does not resume operations or provide responses after the failure. [5]
Omission Failure	The system fails to send or receive necessary messages, impacting communication and task coordination. [21]
Timing Failure	The system's response occurs outside a specified time interval, either too early or too late, causing issues in time-sensitive operations. [21]
Response Failure	The system provides incorrect outputs or deviates from expected state transitions, potentially leading to erroneous results. [5]
Arbitrary Failure	The system produces random or unpredictable responses at arbitrary times, potentially with incorrect or nonsensical data. This type of failure is challenging to diagnose and manage. [5]

Table 2.1: Brief Description of Failure Types

2.2.2 Strategies

Various strategies and mechanisms can be applied to a system to achieve fault tolerance, and these must be chosen to suit the specific system type. This dissertation will primarily focus on software fault tolerance strategies, and focused on those suitable for the programming languages bellow presented. Therefore, next it will be shown some strategies that it will serve as a theoretical basis for some of techniques that it will be used.

Retry Mechanism

The retry mechanism is a widely adopted and straightforward technique that involves reattempting a failed operation under the assumption that transient faults may resolve over time [8]. Despite its simplicity, this strategy is highly suitable in many scenarios, particularly when implementing more complex fault-tolerance mechanisms would introduce unnecessary cost in environments with a high likelihood of transient faults. However, it is crucial to recognize that retrying in the case of a permanent error is pointless. Moreover, if the failure is caused by system overload, uncontrolled retries can aggravate the issue. To address these challenges, implementing a maximum retry limit and incorporating strategies such as exponential backoff, where retries are spaced out with increasing delays, becomes essential [6, 14].

This approach operates by attempting the operation a predefined number of times or until a set timeout is reached. If the retries ultimately fail, the system can fall back on alternative measures, such as logging the failure, invoking a fallback operation, or redirecting the request to another asset [21]. These measures ensure that in the event of a persistent fault, the system will make controlled attempts and roll back in a safe manner.


The retry mechanism's simplicity and low implementation overhead make it ideal for scenarios where the cost of a retry is negligible compared to the complexity of alternative solutions. It is particularly effective in network communication, where transient issues such as dropped packets or server unavailability often resolve with subsequent attempts. Additionally, it is well-suited for database systems to handle transient locking or deadlock conditions, as well as in microservice architectures, where downstream services may temporarily become unresponsive but recover shortly thereafter [14].

This strategy aligns effectively with Actor Models due to their inherent monitoring capabilities, which detect errors and initiate retries automatically. Frameworks such as Akka with Scala have built-in support for this mechanism [21].

Circuit Breaker Pattern

The Circuit Breaker pattern, inspired by electrical circuits, is designed to prevent the failure of a single subsystem from cascading and compromising an entire system. This pattern tries to maintain the overall system stable by isolating failing components [6]. By actively monitoring the health of operations and selectively blocking problematic ones, circuit breakers act as safeguards against system overload and degradation [23].

Circuit breakers operate in three primary states: Closed, Open, and Half-Open [6, 23]. In the Closed state, illustrated in Figure 2.1 by first request, operations proceed as usual, with all requests passing through the circuit breaker while it monitors for potential failures. When failures exceed a predefined threshold within a specified time window, whether measured as a count or a percentage of failed attempts, the circuit breaker transitions to the Open state, illustrated in Figure 2.1 by third request. In this state, all requests are blocked to prevent

Figure 2.1: Diagram illustrating the states of a Circuit Breaker ¹.

greater pressure on the failing subsystem. During this time, it is essential to issue an alert to monitoring systems to ensure operational visibility [6, 23]. After a cool-down period, the circuit breaker moves to the Half-Open state, where it permits a limited number of test requests to verify if the underlying issue has been resolved. If these test requests succeed, the circuit breaker resets to the Closed state and resumes normal operation. Otherwise, it reverts to the Open state [23].

The Circuit Breaker pattern is particularly well-suited for distributed systems, such as microservice architectures, where dependencies on external services can lead to cascading failures [23]. For instance, if a downstream service becomes unresponsive, the circuit breaker blocks further requests, providing the service with time to heal and avoiding the risk of overloading it with retries [6]. It is equally effective in scenarios involving third-party APIs, where temporary rate limits or outages can impact availability. In database systems, circuit breakers can mitigate the effects of resource contention or extended downtime by isolating problematic queries, ensuring the broader system remains operational.

When compared to pure retry mechanisms, the Circuit Breaker pattern provides a more sophisticated approach to fault tolerance. While retries focus on recovering from transient faults, they can harm even more issues under conditions such as system overload or persistent failures [6]. In contrast, circuit breakers proactively block failing operations, reducing the risk of cascading failures and preserving overall system stability. This type of isolation makes circuit breakers a valuable strategy in building fault-tolerant systems.

Replication and Redundancy

Replication is a fundamental strategy for achieving fault tolerance in distributed systems and is widely used across various domains [7]. By creating multiple copies, also called

¹Adapted from Oscar Blancarte Blog. <https://www.oscarblancarteblog.com/2018/12/04/circuit-breaker-pattern/> (accessed 8 December 2024).

replicas, of data or processes, replication eliminates single points of failure, ensuring system reliability, availability, and transparency [9]. This approach allows a system to tolerate faults by introducing redundancy, which distributes operations across a group of replicas rather than relying on a single vulnerable node [5].

To effectively coordinate replicas and maintain consistency, replication mechanisms employ various strategies, which can be categorized as follows [21]:

- **Active Replication (Semi-Active):** In this strategy, all replicas process incoming requests simultaneously, and the system relies on consensus algorithms to maintain consistency among the results.
- **Passive Replication (Semi-Passive):** One replica, designated as the leader or primary, handles all client requests and updates other replicas (backups) with state information. In case of a primary failure, backup replicas are promoted or synchronized to restore the system's functioning.
- **Passive Backup (Fully Passive):** Replicas act as standby backups in this approach. A backup replica is only activated when the primary fails, minimizing overhead during normal operation.

These replication strategies align with the principles of the CAP theorem, which states that a distributed system can guarantee at most two of the following three properties: consistency, availability, and partition tolerance. Replication strategies often emphasize availability and partition tolerance, potentially compromising consistency due to the inherent challenges of achieving consensus and synchronizing data across replicas. Nonetheless, this trade-off enables systems to scale, increase availability, and provide geographical transparency to end users [14].

Consensus Algorithms

Achieving consensus is essential in distributed systems to ensure that a group of processes operates cohesively as a single entity [5]. Consensus algorithms enable replicas to agree on a shared state or a sequence of operations, even in the presence of faults. Two famous used consensus algorithms in distributed systems are Raft and Paxos [5].

Paxos

Paxos first appeared in 1989 and has evolved over time, earning a reputation as a complex and difficult to understand algorithm [5]. Due to its challenges and the availability of newer alternatives, such as Raft, that is described below. This Paxos explanation will focus on its core concepts without delving into exhaustive details.

Paxos ensures that a group of distributed replicas agrees on a single value, even in the presence of faults. It operates under challenging conditions: replicas may crash and recover, messages can be delayed, reordered, or lost, and no assumptions are made about message delivery timing [5, 24]. The algorithm revolves around three distinct roles [9, 24]:

- **Proposers:** Suggest values for the system to agree upon.
- **Acceptors:** Vote on proposals, ensuring fault tolerance by requiring a majority for decisions.
- **Learners:** Observe the final agreed value and disseminate the result across the system.

With the roles defined, the Paxos algorithm progresses through the following phases to achieve consensus [5, 9, 24]:

- **Prepare Phase:** A proposer generates a proposal with a unique sequence number and sends a *prepare request* to a list of acceptors.
 - Acceptors respond with a *promise* not to accept earlier proposals.
 - If an acceptor has already accepted a value, it shares this value with the proposer.
- **Accept Phase:** Based on responses from the prepare phase, the proposer sends an *accept request* with a value:
 - If an acceptor had previously accepted a value, the proposer adopts that value.
 - Otherwise, the proposer chooses its own value.
 - Acceptors respond by accepting the value if it doesn't conflict with their earlier promises.
- **Commit (or Learn) Phase:** Once a majority of acceptors accepts the value, consensus is achieved.
 - The proposer informs all replicas, which then commit the value.

While Paxos is robust in theory and guarantees consistency, its complexity and subtle behaviors make it difficult to implement correctly or faithfully to its original design [5]. Over time, new variations and extensions, such as Multi-Paxos, have been developed. Multi-Paxos enables the system to achieve consensus on multiple values, making it more practical for real-world applications, like Chubby lock service of Google [9], but in general it is not considered a highly adopted algorithm.

The inherent complexity of Paxos has also driven the creation of simpler consensus algorithms, such as Raft, which aim to provide the same guarantees while being easier to understand and implement [5, 24].

Raft

Raft is a consensus protocol designed to enable fault-tolerant operation in distributed systems. It ensures that a process will eventually detect if another process has failed and take appropriate corrective action. Raft was developed as a more comprehensible and practical alternative to Paxos, addressing its complexity and promoting clarity [5, 25]. This algorithm fits on semi passive replication strategy.

Each process in Raft maintains a log of operations, which may include both committed and uncommitted entries. The primary goal of Raft is to ensure that these logs remain consistent across all servers, such that committed operations appear in the same order and position in every log [5]. To achieve this, Raft uses a leader-based approach, where one server assumes the role of leader while the remaining servers act as followers. The leader is responsible for determining the sequence of operations and ensuring their consistent replication [6]. The typical number of nodes is five, which theoretically allows for two failures [25].

When a operation request is submitted, the leader appends the operation to its log as a tuple where it contains: the operation to be executed, the current term of the leader, and the index of the operation in the leader's logs [5]. The term is reset every time an election

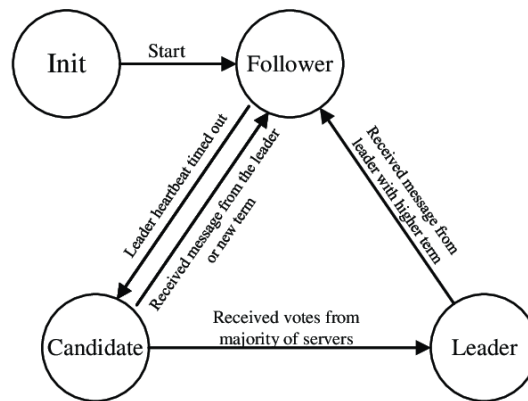


Figure 2.2: Diagram illustrating the states of a node in the Raft algorithm ².

occurs, starting from zero [25]. This information is then propagated to the followers using a process inspired by the two-phase commit protocol, where it consists on [5, 6]:

1. **Append Phase:** The leader sends the new log entry to all followers. The followers append the entry to their logs and send an acknowledgment signal back to the leader.
2. **Commit Phase:** Upon receiving the acknowledgments from a majority of followers, the leader marks the entry as committed, executes the operation, updates its state, and notifies the client of the result. At the same time, the leader informs all followers of the commitment, ensuring their logs reflect the updated status.

This two-step process guarantees that committed entries are replicated on a majority of servers, preserving durability and consistency, even in case of server failures [5]. However, there are cases where the leader fails and an election starts among the followers. The followers acknowledge the leader's failure through the heartbeat strategy, where after a certain time without receiving any signal sent by the leader, the follower starts an election [6, 25], like represented on Figure 2.2 by the change of state from follower to candidate. To prevent multiple followers from initiating elections simultaneously, heartbeat timeouts are randomized. [5, 6].

The change of state of the node is displayed on the Figure 2.2 and the process consists on the following steps [6, 25]:

1. **Transition to Candidate:** A follower transitions to a candidate state, increments its term number, and broadcasts requests for votes from other servers.
2. **Voting:** Each server can vote for one candidate per term. A server grants its vote only if the candidate's log is at least as complete as its own, ensuring that the elected leader has the most up-to-date log.
3. **Leader Selection:** If a candidate receives votes from a majority of servers, it becomes the leader for the current term.

Once elected, the new leader reconciles any inconsistencies by broadcasting missing log entries to followers during subsequent operations, ensuring consistency across the cluster [5].

²Adapted from Jinjie Xu and colleagues. <https://doi.org/10.3390/sym14061122/> (accessed 8 December 2024).

Raft's structured and modular design prioritizes simplicity, reliability, and fault tolerance. Its leader-based model centralizes decision-making and log synchronization, while its robust mechanisms for log replication and leader election ensure consistency and availability even in the face of failures.

Comparison on Paxos and Raft

Both Paxos and Raft are leader-based distributed consensus algorithms designed to ensure consistency in replicated state machines [24]. The key difference lies in their approach to leader election and log replication. Raft prioritizes simplicity and readability, requiring candidates to have up-to-date logs before becoming leaders, avoiding complex log synchronization steps common in Paxos. Paxos, while general and flexible, allows for greater complexity, such as out-of-order log replication and term adjustments during leadership changes [5].

Raft is widely used in distributed systems requiring high availability and consistent state management. Notable applications include etcd for distributed key-value storage, Consul for service discovery, and CockroachDB for scalable, consistent databases [6, 24]. Paxos, on the other hand, is often found in legacy systems and specialized applications like Google's Chubby service [24].

Checkpointing and Message Logging

This strategies focuses more on the recovery part of the faults, so they actuate over an error. This strategies have the goal to recover a error-free state [5].

The checkpointing strategy periodically saves the state of a process so that, in the event of a failure, the process can restart from the last saved state, or "checkpoint," rather than start all over. This approach reduces the need to repeat the entire operation [9, 21].

Message logging is a lighter-weight approach with a similar goal. Instead of saving entire checkpoints, it records all the necessary messages that lead the process to a specific state. In case of a failure, the messages are replayed in the same order, guiding the system back to the desired state [8].

2.3 Distributed and Concurrency Programming

Distributed and concurrent programming languages play an important role in building resilient and fault-tolerant systems [26]. In distributed systems, where components operate across multiple nodes, and in concurrent systems, where tasks can execute in parallel or concurrently on the same machine's Central Processing Unit (CPU), programming languages must provide mechanisms to manage faults effectively. These mechanisms should isolate faults to prevent cascading failures, at the same time ensuring overall system reliability and availability [27], or should have forms to equip the language with capacities to handle this type of systems by frameworks or libraries.

The evolution of distributed programming languages help to address the complexities of developing distributed systems, which include issues such as concurrency, parallelism, fault tolerance, and secure communication [26]. This has driven the evolution of new paradigms, languages, frameworks, and libraries aimed at reducing development complexity in distributed and concurrent systems [12].

2.3.1 Models and Paradigms

The field of distributed programming has been shaped by research and development in concurrency and parallelism, and some models and paradigms have been developed to address this challenge. Some ideas had some focus restricted to the research others have been addressed to the industry. In the following it will be described the models and paradigms that bring interest to this dissertation:

Actor Model

The Actor Model, a conceptual framework for concurrent and distributed computing, was introduced by Carl Hewitt in 1973 [28]. It defines a communication paradigm where an actor, the fundamental unit of computation, interacts with other actors exclusively through asynchronous message passing, with messages serving as the basic unit of communication [29]. Each actor is equipped with its own mailbox, which receives messages and processes them sequentially [30].

A core principle of the Actor Model is isolation, maintaining their own internal state that is inaccessible and immutable by others [30]. This eliminates the need for shared memory, reducing complexity and potential data races [12].

The Actor Model also introduces the concept of supervision, where actors can monitor the behavior of other actors and take corrective actions in the event of a failure. This supervisory mechanism significantly enhances fault tolerance, enabling systems to recover gracefully from errors without compromising overall reliability [29].

The Actor Model has been instrumental in shaping distributed system design and has been natively implemented in programming languages such as Erlang, Clojure and Elixir [31]. Additionally, the model has been extended to other languages through frameworks and libraries. For instance, Akka brings actor-based concurrency to Scala, C# and F# while Kilim offers similar functionality for Java [29]. Comparable patterns can also be adopted in other languages like Go, Rust, and Ruby using libraries or custom abstractions.

Communicating Sequential Processes

The field of distributed computing emphasizes mathematical rigor in algorithm analysis, with one of the most influential models being Communicating Sequential Processes (CSP), introduced by C.A.R. Hoare in 1978 [32].

CSP offers an abstract and formal framework for modeling interactions between concurrent processes through channels, which serve as the communication medium between them [33]. Processes operate independently, but they are coupled via these channels, and communication is typically synchronous, requiring the sender and receiver to synchronize for message transfer [32]. While similar in some respects to the Actor Model, CSP distinguishes itself through its emphasis on direct coupling via channels and synchronization.

The CSP model influenced on programming languages and frameworks. For example, Go integrates CSP concepts in its implementation of goroutines and channels [12, 33, 34]. In addition, the language Occam attempts to offer a direct implementation of CSP principles with its focus on critical projects such as satellites [35].

Microservices Architectures

A significant evolution in designing distributed systems has emerged with the appearance of microservices architectures. This paradigm elevates the focus to a higher level of abstraction, enabling language-agnostic systems by decomposing a monolithic application into a collection of loosely coupled, independently deployable services, each responsible for a specific function [36]. These services communicate using lightweight protocols such as Hypertext Transfer Protocol (HTTP), Google Remote Procedure Call (gRPC), or message queues, fostering separation of concerns, modularity, scalability, and fault tolerance [36].

Microservices architectures allow general purpose programming languages to participate in distributed computing paradigms by leveraging frameworks, libraries, and microservices principles [37].

Although microservices are often associated with strict business principles, their abstract concepts can be adapted to focus on architectural designs that leverage communication middleware for distributed communication. By adopting these principles, it becomes possible to create distributed systems with fault-tolerant capabilities using general-purpose programming languages.

2.3.2 Distributed and Concurrent Programming Languages

Distributed and concurrent programming languages are designed to handle multiple tasks simultaneously across systems or threads. Some languages, such as Java, Rust, and lower-level languages like C with PThreads, require developers to explicitly manage concurrency [12, 33]. These approaches often introduce complexity, increasing the probability of deadlocks or race conditions. This has driven the need for languages and frameworks that abstract away these challenges, offering safer and more developer-friendly concurrency models [12].

One widely adopted paradigm for mitigating concurrency issues is the Actor Model. By avoiding shared state and using message passing for communication, the Actor Model reduces risks inherent in traditional concurrency mechanisms such as mutexes and locks [12]. Erlang, for instance, is renowned for its fault tolerance and “let-it-crash” philosophy, which delegates error handling to its virtual machine [27]. Supervising actors monitor and recover from failures, making Erlang highly suitable for building robust distributed systems [26]. Building on Erlang’s foundation, Elixir introduces modern syntax and developer tooling while retaining Erlang’s strengths for creating large-scale, fault-tolerant systems. These features make Elixir a popular choice for modern distributed systems development [38].

Haskell, a pure functional programming language, provides a deterministic approach to concurrency, ensuring consistent results regardless of execution order [12]. Its extension, Cloud Haskell³, builds upon the Actor Model, drawing inspiration from Erlang, to allow distributed computation through message passing.

Similarly, Akka, a framework built with Scala, adopts the Actor Model to support distributed and concurrent applications. Akka combines Scala’s strengths in functional and object-oriented programming, enabling developers to merge these paradigms effectively [12]. Unlike Erlang, Akka operates on the Java Virtual Machine (JVM), providing seamless interoperability with Java-based systems [39].

³Official website of Cloud Haskell: <https://haskell-distributed.github.io/> (accessed 25 November 2024)

Go, developed by Google, simplifies concurrent programming through its lightweight goroutines and channels, inspired by the CSP paradigm, which abstracts threading complexities [35]. Go's emphasis on simplicity and performance has made it a preferred choice for developing scalable microservices and cloud-native applications, particularly as microservices architectures continue to gain popularity [34].

For specialized use cases like Big Data processing, frameworks such as Hadoop provide distributed computing capabilities tailored to data-intensive tasks. Hadoop abstracts the complexities of handling distributed storage and processing, offering features such as scalability, fault tolerance, and data replication [40].

Other pioneer languages, such as Emerald, Oz, and Hermes, still exist but have minimal community and industry support, as reflected in popularity rankings like RedMonk January 2024⁴ and Tiobe November 2024⁵.

Conversely, some relatively recent languages have gained attention. Unison⁶ employs content-addressed programming using hash references to improve code management and distribution. Gleam⁷ compiles to Erlang and offers its own type-safe implementation of Open Telecom Platform (OTP), Erlang's actor framework. Pony⁸, an object-oriented language based on the Actor Model, introduces reference capabilities to ensure concurrency safety. However, these languages have yet to achieve significant industry adoption, as evidenced by their absence from the RedMonk January 2024 and Tiobe November 2024 rankings.

In Table 2.2, the most relevant languages and frameworks for this theme are presented to facilitate a concise analysis. Additionally, rankings from TIOBE November 2024 and IEEE Spectrum August 2024⁹ are included to provide an overview of their popularity and adoption.

Analyses and Language Choice Justification

The focus of this dissertation is on Elixir as the central language for comparison. Elixir is chosen due to its modern syntax, developer-friendly tooling, and robust foundation on the BEAM virtual machine [38]. Since Elixir inherits all the strengths of Erlang [12], including fault tolerance and the Actor Model, a direct comparison with Erlang is unnecessary as they share the same core runtime and strategies. Such a comparison would likely yield redundant results and add little value to the research.

On the other hand, comparing Elixir with low-level languages like Java, Rust, and C would also be less effective. These languages require explicit management of concurrency and fault tolerance [12], introducing complexities that diverge significantly from Elixir's high-level abstractions. A comparison in this context might be unfair and would not provide meaningful insights given the focus on fault tolerance and distributed systems.

Instead, a comparison with Scala and Akka provides a more relevant perspective. Both Elixir and Akka share the paradigm Actor Model for concurrency and fault tolerance, but their underlying virtual machines differ: the BEAM for Elixir and the JVM for Akka [39].

⁴RedMonk January 2024: <https://redmonk.com/sograde/2024/03/08/language-rankings-1-24/> (accessed 28 November 2024)

⁵Tiobe November 2024: <https://www.tiobe.com/tiobe-index/> (accessed 28 November 2024)

⁶Official website of Unison: <https://www.unisonweb.org/> (accessed 27 November 2024)

⁷Official website of Gleam: <https://gleam.run/> (accessed 27 November 2024)

⁸Official website of Pony: <https://www.ponylang.io/> (accessed 27 November 2024)

⁹IEEE Spectrum 2024: <https://spectrum.ieee.org/top-programming-languages-2024/> (accessed 28 November 2024)

Name	Concurrency Strategy	Model	TIOBE Nov 2024	IEEE Spectrum 2024
Java	Explicit	Object-Oriented	3	2
Rust	Explicit	Procedural	14	11
C (PThreads)	Explicit	Procedural	4	9
Erlang	Actor Model	Functional	50+	48
Elixir	Actor Model	Functional	44	35
Haskell	Evaluation Strategy	Functional	34	38
Scala (Akka)	Actor Model	Functional	30	16
Go	CSP	Procedural	7	8
Hadoop	Distributed Framework	Procedural	N/A	N/A
Unison	Hash References	Functional	N/A	N/A
Gleam	Actor Model	Functional	N/A	N/A
Pony	Actor Model	Object-Oriented	N/A	N/A

Table 2.2: Characteristics of Distributed and Concurrent Programming Languages

Additionally, Scala with Akka is notable for its community acceptance [12]. This comparison is valuable because it explores how different implementations of the same paradigm can influence fault tolerance strategies and performance.

Furthermore, too recent or older languages with minimal popularity, such as Emerald, Oz, Unison and Gleam are excluded from this study. These languages lack widespread adoption, and insights derived from them would have limited applicability for the majority of developers, as demonstrated in Table 2.2 with a non-appearance in the Tiobe and IEEE Spectrum rankings.

From another perspective, the inclusion of Go in this study adds an interesting dimension to the comparison. Go, unlike Elixir and Akka, does not have built-in support for native distributed systems. However, its increasing popularity and industry adoption make it a strong candidate for exploration [35]. By examining how Go can achieve fault tolerance using libraries and abstractions under a microservices strategy, the study can assess whether an external abstraction layer can match or exceed the capabilities of languages with native support. This investigation could reveal whether the flexibility of a non-native distributed model can compensate for the lack of built-in features, providing valuable insights for developers operating in modern cloud-native environments.

Chapter 3

Literature Review

3.1 Research Questions

- RQ1: How do the programming languages Elixir, Scala/Akka, and Go implement fault-tolerance mechanisms in distributed systems, and what are the comparative strengths, weaknesses, and trade-offs of each approach?
- RQ2: What are the most effective benchmarking strategies for simulating distributed environments, and which metrics are most relevant for evaluating fault tolerance in these scenarios?

3.2 State of Art

3.3 Conclusions

Chapter 4

Planning

4.1 Project Charter

The project charter provides a overview of the stakeholders, benefits, and assumptions.

Stakeholder

Identification	Power	Interest
Students or developers in the areas of distributed systems and fault tolerance	Low	Low
Administration of the master's program, responsible for dissertation evaluation	Medium	Low
Advisor	High	High

The administration of the master's program has a moderate influence due to institutional requirements, but a potential interest in the results, since it will be a project for the institution. The advisor, on the other hand, has a high level of influence and interest, as his guidance and approval are fundamental to the success of the project. In addition, students and developers in general may have an interest in the outcome in order to have a guide for a better choice of their projects, but low influence.

Benefits

- **Decision Support for Developers:** The project will provide a detailed analysis of fault tolerance aspects in Elixir, Go, and Scala with Akka, offering developers and system architects a practical guide to help them choose the most suitable language for specific fault-tolerant distributed systems scenarios.
- **Open Source Opportunities:** The findings could reveal areas for improvement in the evaluated languages, inspiring open-source developers to create libraries, frameworks, or enhancements to already existing ones.
- **Academic Contributions:** The dissertation will contribute to the existing body of knowledge in the areas of distributed systems, fault tolerance, and microservices. It will provide insights into comparative aspects in the languages of debate.

Assumptions

- **Computational Resources:** It is assumed that the available computational resources, including hardware and software tools, will suffice to simulate the benchmarking scenarios for each language under realistic system conditions.




Figure 4.1: The WBS of the project.

- **Support and Guidance:** The advisor will provide timely and effective feedback on each deliverable, ensuring alignment with project objectives.
- **Consistency Across Languages:** The chosen languages (Elixir, Go, Scala with Akka) have sufficient community support, libraries, and tools to implement the required benchmarking scenarios consistently.

4.2 Work Breakdown Structure

The objective of the Work Breakdown Structure (WBS) is to detail the project scope in a visual and hierarchical manner, enabling a clear understanding of how each deliverable connects to the overall project. As shown in Figure 4.1, the main point of this project is the dissertation. With the objective defined, the first phase focuses on project planning. This phase establishes the foundation by outlining the project charter, creating a WBS, and developing a timeline through a Gantt chart.

Once the planning phase is complete, the subsequent phases align with the Design and Creation research method. This research method was chosen given the nature of the project, because while the final objective is clear, there are uncertainties about how to achieve each stage, as every step builds on the outcomes of the previous one. Consequently, the method divides the project into sequential phases: design, implementation, and conclusion. Each phase has clearly defined deliverables that align with the WBS, ensuring that progress can be monitored and adjustments can be made as needed.

The final phase, the conclusion, consolidates all findings and results, translating them into the completed dissertation.

Work Breakdown Structure Dictionary. Following is described the WBS Dictionary that as the responsibility of detailing each phase in order to be defined what are the goals and the acceptance criteria in a concise and clear way.

Item Name	Type of Item	Additional Description / Acceptance Criteria
(1.1) Planning	Phase	This phase includes all initial project setup tasks.

Item Name	Type of Item	Additional Description / Acceptance Criteria
(1.1.1) Project Charter	Deliverable	The project charter must be created following the project's scope and management guidelines. Acceptance Criteria: The project charter must be approved by the advisor.
(1.1.2) WBS	Deliverable	The WBS should break down the project into manageable components. Acceptance Criteria: The WBS should be validated by the advisor and include all project elements.
(1.1.3) Gantt Chart	Deliverable	A detailed timeline outlining tasks, dependencies, competence development plan, milestones, and the dissertation deadline. Acceptance Criteria: The Gantt chart must accurately reflect project phases and be reviewed by the advisor.
(1.2) Research	Phase	This phase focuses on gathering the required knowledge and literature to support the project.
(1.2.1) Objectives	Deliverable	Clear objectives for the project, that must detail what are the expected outcomes. Acceptance Criteria: Objectives should align with the research goals and be validated by the advisor.
(1.2.2) Background	Deliverable	Research and summarize the background of fault tolerance in distributed systems and the distributed and concurrent programming languages. Acceptance Criteria: The background section should include sufficient theoretical content approved by the advisor, and must include a clear justification for the languages chosen.
(1.2.3) State-of-art	Deliverable	Review the current literature on fault tolerance in Elixir, Go, and Scala with Akka. Also, what are the latest techniques for benchmarking distributed and concurrent programming, and if there are already studies on this topic. Acceptance Criteria: State-of-the-art review must highlight gaps and relevance to the project scope.
(1.3) Design	Phase	This phase involves requirements gathering, testing plan, and benchmarking design.
(1.3.1) Requirements Gathering	Deliverable	Collect requirements for the benchmarking and evaluation of fault tolerance aspects. Acceptance Criteria: Requirements must be detailed, reviewed, and approved by the advisor.

Item Name	Type of Item	Additional Description / Acceptance Criteria
(1.3.2) Testing Plan	Deliverable	A plan for testing different fault tolerance strategies and mechanisms in Elixir, Go, and Scala with Akka. Acceptance Criteria: Testing plan must include scenarios and validation methods, reviewed by the advisor.
(1.3.3) Benchmarking Design	Deliverable	Define the design for benchmarking environments. Acceptance Criteria: Benchmarking environments design must be validated by the advisor, and must adhere to the test plan created.
(1.4) Implementation	Phase	This phase involves the development of benchmarking prototypes.
(1.4.1) Prototypes	Deliverable	Develop prototypes in Elixir, Go, and Scala with Akka for fault tolerance testing. Acceptance Criteria: Prototypes must meet the test plan previously created and must be supported on the benchmarking design planned.
(1.4.2) Test Triggers	Deliverable	Create fault injection mechanisms for testing fault tolerance. Acceptance Criteria: Fault injection methods must simulate real-world scenarios and be validated by tests.
(1.4.3) Observability	Deliverable	Implement observability tools for monitoring system behavior during tests. Acceptance Criteria: Observability setup must capture the metrics defined on the test validations methods.
(1.5) Evaluation	Phase	Evaluate the results of the benchmarking tests.
(1.5.1) Benchmarking Result	Deliverable	Analyze and document the outcomes of benchmarking fault tolerance aspects. Acceptance Criteria: Results must be clear, reproducible, and reviewed by the advisor.
(1.6) Conclusions	Phase	Finalize and present the results of the dissertation.
(1.6.1) Dissertation	Deliverable	Compile the dissertation document with findings and analyses. Acceptance Criteria: Dissertation must meet academic formatting and content guidelines.
(1.6.2) Conclusions	Deliverable	Write concise conclusions summarizing key findings from the research, with the goal of creating a guide for future developers consult. Acceptance Criteria: Conclusions must be concise and detail what are the cons and pros of using each language for each specific case, so that developers can easily decide.

Item Name	Type of Item	Additional Description / Acceptance Criteria
(1.6.3) Presentation	Deliverable	Prepare and deliver the final presentation to the evaluation committee. Acceptance Criteria: Presentation must be clear and precise.

Table 4.1: WBS dictionary

4.3 Gantt Diagram

4.3.1 Project Management and Scheduling

For project management, the start date was established as 02/02/2025, after the 1^o semester exams. The scheduling mode was set to automatically, where it calculates the dates based on the days of each task. The schedule is based on a calendar with no restrictions on working days, meaning all days, including weekends, are considered working days to ensure continuous progress.

melhorar esta parte de baixo

Estimation Rationale. Task duration estimates were determined based on the complexity and importance of each phase. The background and state-of-the-art sections were assigned the highest weights due to their critical role in forming the foundation of the research. These sections require extensive literature review and analysis, which are time-intensive. After completing the research phase, the design stage was given significant weight, as it involves defining requirements and developing a comprehensive testing plan, both of which are essential for a successful implementation.

The duration of the implementation phase was estimated considering the iterative nature of prototyping and testing triggers, ensuring sufficient time for observability and adjustments. The evaluation phase, particularly benchmarking results, was scheduled with adequate time to ensure thorough analysis and interpretation of the findings. Lastly, the conclusions phase was planned with sufficient time to synthesize the project outcomes and finalize the dissertation.

Resource Allocation and Monitoring. A resource table was included in the project plan to track resource allocation, focusing on the student and advisor as the primary stakeholders. The advisor's role is focused on controlling and monitoring tasks, providing feedback and guidance throughout the project.

Cost Management. The cost component was excluded from the project plan, as it is not applicable to this type of academic project.

4.3.2 Monitoring and Controlling Procedures

The strategy defined to have control over the progress of each task was to mark its progress by an estimated percentage of completion. This strategy ensures that the progress of the task is clearly visible and measurable, giving both the student and the advisor a view of progress, like it is possible to observe the “% Completed” column on the Figure 4.2.

To improve this process, specific tasks dedicated to monitoring and control have been incorporated, such as “*Validation and Refinement with the Advisor*”, like illustrated on Figure 4.2.




Figure 4.2: Monitoring and control procedures displayed on the Gantt chart.

These tasks have two main objectives, firstly, to make the student responsible for presenting the partial document ready to be assessed. Secondly, to ensure that the advisor is warned in advance of the need for closer and more active feedback on these pre-defined days. Feedback can be given asynchronously via messages or synchronously during scheduled meetings. After the feedback time is up, the goal is to be able to present a refined partial document, referred to as the *“Document version X”* task, also possible to observe on Figure 4.2.

Additionally, milestones have been defined to mark the completion of each significant project phase. These milestones serve as checkpoints to ensure progress is on track and can be observed in Figure 4.2 under the main task *“Milestones”*.

To manage potential delays, a baseline has been established. This baseline records the initial schedule, allowing deviations to be tracked throughout the project. This mechanism provides an overview of delays and their impact on the schedule. The *“Variance”* column in Figure 4.2 illustrates this feature, allowing a visualization of changes between planned and actual progress.

4.3.3 Meeting Sessions

To ensure consistent communication and effective progress monitoring with the advisor, a series of biweekly meetings has been scheduled on Wednesdays, with each session expected to last between 30 minutes and 1 hour. While the schedule includes a predefined list of sessions, it remains flexible, allowing adjustments to the frequency of meetings as the project evolves. For instance, the number of meetings may increase during the final stages of the project, at which point the Gantt chart will be updated accordingly.

The meeting schedule is illustrated in Figure 4.3, which includes 11 recurring tasks organized under the main task *“Meeting Sessions”*. These meetings are planned to take place online.

4.3.4 Competencies Development Plan

To address the competencies identified during the diagnosis of critical skills required for the completion of the dissertation, a dedicated section titled *“Competencies Development Plan”* has been incorporated into the project plan, as illustrated in Figure 4.4. This section outlines targeted tasks designed to address these areas for improvement.




Figure 4.3: Meeting sessions represented on the Gantt chart.

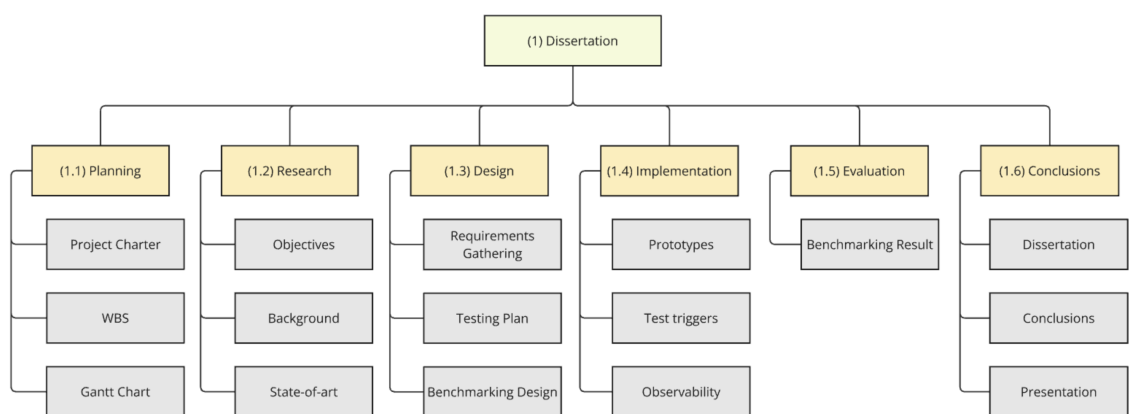


Figure 4.4: Competencies Development Plan represented on the Gantt chart.

For stress management and resilience, the plan includes attending a Stress Management course.¹ To enhance self-discipline and time management, the initial task involves identifying and installing at least one productivity application, such as tools that restrict smartphone usage during certain periods, for example. Additionally, time management competence is further explored by taking the Time Management Fundamentals course².

Communication, another key area of focus, will be developed by attending to the "Communicating with Confidence" course³.

4.4 Risk Management

Effective risk management seeks to transform potential uncertainties into more predictable and controlled outcomes. To achieve this, the most significant risks identified are described next

4.4.1 Risk 1: Bugs in Third-Party Libraries

Description: There is a potential risk of encountering bugs in third-party libraries, which could compromise the viability of implementation and testing of the prototypes. Since the project relies on external libraries to implement fault-tolerant strategies, the stability and reliability of these libraries are important.

Cause: The cause of this risk is the need of trust on external software components.

Effect on the Project: Errors in the libraries can compromise the viability of the prototype's development and also the integrity of the results.

Risk Owner: Student.

Probability. 2. **Impact:** 4. **PI Score.** 8.

Risk Response Strategy: To mitigate this risk, alternative libraries will be identified for each strategy and language in advance. At least two or three libraries will be shortlisted and prioritized. If the primary library encounters significant bugs or issues, the next library on the list will be utilized.

Expected Result Without Action: If no action is taken, delays in prototype development will occur.

Risk Response Type: Mitigation.

Response Description: A proactive approach will be taken to evaluate multiple libraries during the research phase.

¹Managing Stress Course: <https://www.linkedin.com/learning/managing-stress-2019/> (accessed 4 December 2024)

²Time Management Fundamentals Course: <https://www.linkedin.com/learning/time-management-fundamentals-14548057/the-power-of-managing-your-time/> (accessed 4 December 2024)

³Communicating with Confidence Course: <https://www.linkedin.com/learning/como-se-comunicar-com-confianca/> (accessed 4 December 2024)

4.4.2 Risk 2: Integration Challenges Between Components

Description: Integration issues could arise when combining multiple components, such as third-party libraries, testing frameworks, and custom implementations.

Cause: Differences in interfaces, versions, or dependencies among components used in the project.

Effect on the Project: These challenges could delay testing and result in compatibility issues that reduce productivity.

Risk Owner: Student.

Probability. 2. **Impact:** 4. **PI Score.** 8.

Risk Response Strategy: To mitigate this risk, dependencies and versions will be carefully managed using dependency management tools (e.g., *mix* for Elixir, *go.mod* for Go, *sbt* for Scala). Component integration will also be tested incrementally to identify issues early.

Expected Result Without Action: Significant delays during the integration phase.

Risk Response Type: Mitigation.

Response Description: Incremental integration and version control practices will ensure smoother component interaction.

Bibliography

- [1] Dipankar Deb, Rajeeb Dey, and Valentina E. Balas. *Engineering Research Methodology*. Dec. 2018. doi: 10.1007/978-981-13-2947-0. url: <https://doi.org/10.1007/978-981-13-2947-0>.
- [2] NSPE. *NSPE Code of Ethics for Engineers*. Accessed at 01.12.2024. url: <https://www.nspe.org/resources/ethics/code-ethics>.
- [3] IEEE. *IEEE Code of Ethics*. Accessed at 01.12.2024. url: <https://www.ieee.org/about/corporate/governance/p7-8.html>.
- [4] ACM. *ACM Code of Ethics and Professional Conduct*. Accessed at 01.12.2024. url: <https://www.acm.org/code-of-ethics/>.
- [5] S. Andrew Tanenbaum and M. Maarten Van Steen. *Distributed systems*. 4th ed. Maarten Van Steen, 2023. isbn: 978-90-815406-3-6.
- [6] Roberto Vitillo. *Understanding Distributed Systems: What every developer should know about large distributed applications*. 2021. isbn: 1838430202.
- [7] Arif Sari and Murat Akkaya. "Fault Tolerance Mechanisms in Distributed Systems". In: *International Journal of Communications, Network and System Sciences* 08 (12 2015), pp. 471–482. issn: 1913-3715. doi: 10.4236/ijcns.2015.812042.
- [8] Mohamed. Amroune, Makhlof. Dourdour, and Ahmed. Ahmim. *Fault Tolerance in Distributed Systems: A Survey*. IEEE, 2018. isbn: 9781538642382.
- [9] George Coulouris et al. *Distributed Systems - Concepts and Design*. 2012. isbn: 978-0-13-214301-1.
- [10] Waseem Ahmed and Yong Wei Wu. "A survey on reliability in distributed systems". In: *Journal of Computer and System Sciences*. Vol. 79. Academic Press Inc., 2013, pp. 1243–1255. doi: 10.1016/j.jcss.2013.02.006.
- [11] Atlassian. *Reliability vs availability: Understanding the differences*. Accessed at 16.10.2024. url: <https://www.atlassian.com/incident-management/kpis/reliability-vs-availability>.
- [12] Ivan Valkov, Natalia Chechina, and Phil Trinder. "Comparing languages for engineering server software: Erlang, go, and scala with akka". In: *Proceedings of the ACM Symposium on Applied Computing*. Association for Computing Machinery, Apr. 2018, pp. 218–225. isbn: 9781450351911. doi: 10.1145/3167132.3167144.
- [13] Dominic Lindsay et al. "The evolution of distributed computing systems: from fundamental to new frontiers". In: *Computing* 103 (8 Aug. 2021), pp. 1859–1878. issn: 14365057. doi: 10.1007/s00607-020-00900-y.
- [14] Martin Kleppmann. *Designing Data Intensive Applications*. 2017. isbn: 978-1449373320.
- [15] Nitin Naik. "Performance Evaluation of Distributed Systems in Multiple Clouds using Docker Swarm". In: *15th Annual IEEE International Systems Conference, SysCon 2021 - Proceedings*. Institute of Electrical and Electronics Engineers Inc., Apr. 2021. isbn: 9781665444392. doi: 10.1109/SysCon48628.2021.9447123.
- [16] AWS - Amazon. *Challenges with distributed systems*. Accessed at 03.11.2024. url: <https://aws.amazon.com/builders-library/challenges-with-distributed-systems/>.

- [17] Ivan Beschastnikh et al. "Visualizing Distributed System Executions". In: *ACM Transactions on Software Engineering and Methodology* 29 (2 Mar. 2020). issn: 15577392. doi: 10.1145/3375633.
- [18] IBM. *What is the CAP theorem?* Accessed at 03.11.2024. Aug. 2024. url: <https://www.ibm.com/topics/cap-theorem>.
- [19] Seth Gilbert and Nancy A Lynch. *Perspectives on the CAP Theorem*. 2012. doi: 10.1109/MC.2011.389.
- [20] Ahmad Shukri Mohd Noor, Nur Farhah Mat Zian, and Fatin Nurhanani M. Shai-ful Bahri. "Survey on replication techniques for distributed system". In: *International Journal of Electrical and Computer Engineering* 9 (2 Apr. 2019), pp. 1298–1303. issn: 20888708. doi: 10.11591/ijece.v9i2.pp1298-1303.
- [21] Sucharitha Isukapalli and Satish Narayana Srirama. *A systematic survey on fault-tolerant solutions for distributed data analytics: Taxonomy, comparison, and future directions*. Aug. 2024. doi: 10.1016/j.cosrev.2024.100660.
- [22] Federico Reghenzani, Zhishan Guo, and William Fornaciari. "Software Fault Tolerance in Real-Time Systems: Identifying the Future Research Questions". In: *ACM Computing Surveys* 55 (14 Dec. 2023). issn: 15577341. doi: 10.1145/3589950.
- [23] Martin Fowler. *Circuit Breaker*. 2014. url: <https://martinfowler.com/bliki/CircuitBreaker.html>.
- [24] Heidi Howard and Richard Mortier. "Paxos vs Raft: Have we reached consensus on distributed consensus?" In: *Proceedings of the 7th Workshop on Principles and Practice of Consistency for Distributed Data, PaPoC 2020*. Association for Computing Machinery, Inc, Apr. 2020. isbn: 9781450375245. doi: 10.1145/3380787.3393681.
- [25] Diego Ongaro and John Ousterhout. "In Search of an Understandable Consensus Algorithm". In: *2014 USENIX Annual Technical Conference (USENIX ATC 14)*. Philadelphia, PA: USENIX Association, June 2014, pp. 305–319. isbn: 978-1-931971-10-2. url: <https://www.usenix.org/conference/atc14/technical-sessions/presentation/ongaro>.
- [26] Joe Armstrong. *Early Praise for Programming Erlang, Second Edition*. 2013. isbn: 978-1-937785-53-6.
- [27] Jan Henry Nystrom. *Fault Tolerance in Erlang*. 2009. isbn: 978-91-554-7532-1.
- [28] Carl Hewitt, Peter Bishop, and Richard Steiger. "A universal modular ACTOR formalism for artificial intelligence". In: *Proceedings of the 3rd International Joint Conference on Artificial Intelligence. IJCAI'73*. Stanford, USA: Morgan Kaufmann Publishers Inc., 1973, pp. 235–245. doi: doi/10.5555/1624775.1624804.
- [29] Phil Trinder et al. "Scaling reliably: Improving the scalability of the Erlang distributed actor platform". In: *ACM Transactions on Programming Languages and Systems* 39 (4 Aug. 2017). issn: 15584593. doi: 10.1145/3107937.
- [30] Joeri De Koster, Tom Van Cutsem, and Wolfgang De Meuter. "43 years of actors: a taxonomy of actor models and their key properties". In: *Proceedings of the 6th International Workshop on Programming Based on Actors, Agents, and Decentralized Control. AGERE 2016*. Amsterdam, Netherlands: Association for Computing Machinery, 2016, pp. 31–40. isbn: 9781450346399. doi: 10.1145/3001886.3001890. url: <https://doi.org/10.1145/3001886.3001890>.
- [31] Aidan Randtoul and Phil Trinder. "A reliability benchmark for actor-based server languages". In: *Erlang 2022 - Proceedings of the 21st ACM SIGPLAN International Workshop on Erlang*. Association for Computing Machinery, Inc, Sept. 2022, pp. 21–32. isbn: 9781450394352. doi: 10.1145/3546186.3549928.

- [32] C. A. R. Hoare. "Communicating sequential processes". In: *Commun. ACM* 21.8 (Aug. 1978), pp. 666–677. issn: 0001-0782. doi: 10.1145/359576.359585. url: <https://doi.org/10.1145/359576.359585>.
- [33] Ciprian Paduraru and Marius Constantin Melemciuc. "Parallelism in C++ Using Sequential Communicating Processes". In: *Proceedings - 17th International Symposium on Parallel and Distributed Computing, ISPDC 2018*. Institute of Electrical and Electronics Engineers Inc., Aug. 2018, pp. 157–163. isbn: 9781538653302. doi: 10.1109/ISPDC2018.2018.00030.
- [34] Go. *Official documentation of Go programming language*. Accessed at 10.11.2024. url: <https://go.dev/doc/>.
- [35] Matilde Brolos, Carl Johannes Johnsen, and Kenneth Skovhede. "Occam to Go translator". In: *Proceedings - 2021 Concurrent Processes Architectures and Embedded Systems Conference, COPA 2021*. Institute of Electrical and Electronics Engineers Inc., Apr. 2021. isbn: 9781728166834. doi: 10.1109/COPA51043.2021.9541431.
- [36] Pooyan Jamshidi et al. "Microservices: The Journey So Far and Challenges Ahead". In: *IEEE Software* 35.3 (2018), pp. 24–35. doi: 10.1109/MS.2018.2141039.
- [37] Claudio Guidi et al. "Microservices: A language-based approach". In: Springer International Publishing, Nov. 2017, pp. 217–225. isbn: 9783319674254. doi: 10.1007/978-3-319-67425-4_13.
- [38] Saša Jurić and Francesco Cesarini. *Elixir in Action, Third Edition*. 2024. isbn: 9781633438514.
- [39] Francisco Lopez-Sancho Abraham. *Akka in Action, Second Edition*. Simon and Schuster, 2023. isbn: 9781617299216.
- [40] Ivanilton Polato et al. *A comprehensive view of Hadoop research - A systematic literature review*. 2014. doi: 10.1016/j.jnca.2014.07.022.